

GeoSample U-Net: Teaching & Implementation Guide

Intuition, Mathematical Foundations (Simplified), Architectural Mechanics & Practical Project Results

Pedagogical Focus: Intuition & Concepts	Core Principle: Steer Sampling, Not Grid	Model Capacity: 4,841,497 Params (4.84M)	Efficiency: -38.3% Params, 86.41% Dice
--	---	---	---

1. The Research Paper's Core Intuition: Why GeoSample?

Standard convolutional neural networks suffer from an inherent blind spot: **the rigid Cartesian grid**. A standard 3×3 convolution always samples pixels at fixed rectangular offsets: $[(-1,-1), (-1,0), (-1,1), \dots (1,1)]$. However, biology does not grow on a square grid! In brain MRI scans, tumor boundaries, cortical sulci, and infiltrative lesions are continuous, diagonal, and curved.

When a rigid square kernel sweeps across a curved tumor boundary, it causes two critical problems:

- Boundary Contamination:** The kernel straddles both inside the tumor and outside in healthy brain tissue, blurring the edge.
- Parameter Waste:** To detect boundaries oriented at 15°, 30°, 45°, 60°, etc., a standard CNN is forced to allocate dozens of separate 3×3 filters—memorizing every possible angle through brute-force channel expansion (inflating model size to 8M–30M parameters).

■ THE PARADIGM SHIFT: STEER THE SAMPLING, NOT THE KERNEL GRID

- The Core Insight:** Any boundary point is intrinsically defined by two geometric properties: an **orientation angle (θ)** and a **scale / step size (r)**.
- Instead of sampling a rigid 9-point square, GeoSample places **4 continuous symmetric probes** around each pixel (2 along the edge, 2 across the edge).
- From these 4 symmetric probes, elementary numerical calculus extracts **Context (local average)**, **Gradient (edge sharpness)**, and **Curvature (surface bending)** directly!

The 'Aha!' Moment for Students: Because spatial geometry is already completely computed by the continuous probes and differential tokens, the neural network does NOT need heavy 3×3 spatial convolutions! It only needs lightweight 1×1 pointwise convolutions to blend the physical tokens together. This architectural elegance is why GeoSample cuts model parameters by ~40% while preserving ultra-sharp tumor contours.

2. The Three Core GeoSample Operators (Under the Hood)

The GeoSample framework replaces the three standard U-Net operations with three geometry-guided equivalents:

Standard CNN Block	GeoSample Operator	Parameter Impact	Key Functional Advantage
DoubleConv (Two 3×3 convs)	GeoSample2D	~50% fewer params	4-probe oriented sampling + 6 differential tokens + 1×1 mixing
MaxPool2d(2×2)	GeoDownsample2D	Zero extra params	Dual-path downsampling: AvgPool context + Absolute edge compensation
Naive Skip Cat ([skip, dec])	ConsensusSkip2D	Tiny 1×1 gate	Cosine agreement gating between encoder and decoder boundary angles

Operator Mechanics Breakdown:

- GeoSample2D Execution:** (1) GeometryHead predicts orientation vector $u1=[\cos \theta, \sin \theta]$, orthogonal vector $u2=[-\sin \theta, \cos \theta]$, and step sizes r . (2) Bottleneck projection downsamples channels ($C_{in} // 2$) to prevent memory explosion. (3) Continuous bilinear sampling gathers 4 symmetric neighbors. (4) Arithmetic extraction produces 6 differential tokens [center, context1, context2, Gx, Gy, Curvature]. (5) Two 1×1 convolutions blend tokens to C_{out} . (6) Residual addition.
- GeoDownsample2D:** Solves MaxPool's fatal flaw (which erases opposite edge slopes). Path A performs smooth average pooling for background context. Path B computes gradient magnitudes $|Gx|, |Gy|$ and curvature $|L|$ before pooling, pools them in absolute form, and injects them back as an edge compensation signal.

- **ConsensusSkip2D:** Eliminates boundary ghosting. When the blurry upsampled decoder meets crisp encoder features, it computes the cosine similarity between their predicted boundary angles. High agreement gates the decoder features through; discordant signals are suppressed.

3. Adaptation & Practical Evolution in Our Major Project

While the original 2025 paper proposed GeoSample for 3D volumetric CT/MRI tensors (demanding high-end computing clusters), our project engineered a clean 2D formulation tailored for brain MRI slice segmentation on Kaggle BRISC 2025, trainable within Google Colab's free-tier T4 GPU (16 GB VRAM).

Experimental Run	Base Feats	Parameters	Optimization & Loss	Experimental Finding
Run 1 (Initial)	32	3.52M	Dice+BCE, Plateau LR	85.8% Dice. Proved 2D formulation works; slight underfitting.
Run 2 (Champion)	48	4.84M	Dice+BCE, Cosine LR	86.41% Mean / 93.46% Median Dice. Matches baseline with -38.3% params!
Run 3 (Over-tuned)	60	7.56M	Focal Loss + Deep Supervision	84.9% Dice. Auxiliary heads and focal loss disrupted gradient stability.

■ **TEACHING TAKEAWAY: OCCAM'S RAZOR IN DEEP LEARNING**

- More parameters and complex loss functions do NOT always mean better performance.
- Run 3 proved that adding multi-scale auxiliary heads and Focal Loss on a small dataset (3,540 slices) created conflicting gradient pressures.
- Run 2 (4.84M params, base_features=48, Dice+BCE, Cosine LR) proved that clean geometry-guided sampling is far superior to brute-force parameter scaling.

4. Quantitative Verification: Baseline U-Net vs. GeoSample Run 2

Evaluation on 860 independent test slices under single-pass deterministic inference confirms near-perfect parity:

Evaluation Metric	Baseline U-Net	GeoSample U-Net (Run 2)	Variance / Efficiency Gain
Total Trainable Parameters	7,849,601 (7.85M)	4,841,497 (4.84M)	-38.3% Parameter Savings (~3.0M fewer weights)
Mean Test Dice Score	86.79%	86.41%	-0.38% (Statistically indistinguishable)
Median Test Dice Score	93.86%	93.46%	-0.40% (Identical core overlap robustness)
Mean Test IoU Score	80.00%	79.38%	-0.62% difference
Median Test IoU Score	88.44%	87.74%	-0.70% difference
Excellent Tier (Dice ≥ 0.85)	649 slices (75.5%)	646 slices (75.1%)	Identical proportion of clinical-grade segmentations

5. Empirical Training Curves & Qualitative Visualizations




Figure 1: GeoSample U-Net (Run 2) Training and Validation Curves across 40 epochs.




Figure 2: Qualitative Test Inferences from GeoSample U-Net (Input MRI, Ground Truth, Predicted Mask, and Error Map).

6. Student Viva & Examination Cheatsheet (Quick FAQs)

Q1: Why does GeoSample have 38% fewer parameters than standard U-Net?

Standard U-Net relies on pairs of heavy 3×3 convolutions at every stage to memorize all possible edge orientations. GeoSample replaces these with 4 continuous geometry-steered probes and lightweight 1×1 token mixers. Because spatial geometry is solved by the probes, only channel mixing is needed, slashing weights from 7.85M to 4.84M.

Q2: What are differential tokens and why do they matter?

From 4 symmetric neighbor probes, GeoSample arithmetically extracts Context (local average intensity), Gradient (central difference representing edge jump/direction), and Curvature (second-order difference representing surface bend). These provide explicit physical boundary representations without needing any learnable convolutional parameters.

Q3: What makes GeoDownsample better than MaxPool?

MaxPool2d discards 75% of pixels, often erasing opposite-sign boundary gradients. GeoDownsample uses a dual-path: Path A average-pools context, while Path B computes gradient magnitudes $|G_x|$, $|G_y|$ and curvature $|L|$ before pooling and injects them back as edge compensation.

Q4: What problem does ConsensusSkip solve?

Upsampled decoder features have blurry, slightly misaligned boundaries. Naive concatenation causes 'ghosting' (two competing edge signals). ConsensusSkip computes the cosine similarity between encoder and decoder boundary angles, gating the decoder features only where geometric consensus exists.